

Homework 5: Poisson matrix factorization

STAT 348, UChicago, Spring 2025

Your name here: Haolin Yang

Hours spent: 10

(Please let us know how many hours in total you spent on this assignment so we can calibrate for future assignments. Your feedback is always welcome!)

 [Open in Colab](#)

Instructions

This homework focuses on themes in lectures 10-12 on coordinate ascent variational inference (CAVI), admixture models, and Poisson matrix factorization.

For reference, this homework is a close adaption of [HW5 for Scott Linderman's STATS 305C](#).

Assignment is due **Wednesday May 14, 11:59pm** on GradeScope.

Background

Poisson matrix factorization (PMF) is a mixed membership model like LDA, and it has close ties to non-negative factorization of count matrices. Let $\mathbf{X} \in \mathbb{N}^{N \times M}$ denote a count matrix with entries $x_{n,m}$. We model each entry as a Poisson random variable,

$$x_{n,m} \sim \text{Po}(\boldsymbol{\theta}_n^\top \boldsymbol{\phi}_m) = \text{Po}\left(\sum_{k=1}^K \theta_{n,k} \phi_{m,k}\right),$$

where $\boldsymbol{\theta}_n \in \mathbb{R}_+^K$ and $\boldsymbol{\phi}_m \in \mathbb{R}_+^K$ are *non-negative* feature vectors for row n and column m , respectively.

PMF has been used for recommender systems, aka collaborative filtering. In a recommender system, the rows correspond to users, the columns to items, and the entries $x_{n,m}$ to how much user n liked item m (on a scale of 0, 1, 2, . . . stars, for example). The K feature dimensions capture different aspects of items that users may weight in their ratings.

Note that the Poisson rate must be non-negative. It is sufficient to ensure θ_n and ϕ_m are non-negative. To that end, PMF uses gamma priors,

$$\begin{aligned}\theta_{n,k} &\sim \text{Ga}(\alpha_\theta, \beta_\theta) \\ \phi_{m,k} &\sim \text{Ga}(\alpha_\phi, \beta_\phi),\end{aligned}$$

where α_\star and β_\star are hyperparameters. When $\alpha_\star < 1$, the gamma distribution has a sharp peak at zero and the prior induces sparsity in the feature vectors.

Latent variable formulation

PMF can be rewritten in terms of a latent variable model. Note that,

$$\begin{aligned}x_{n,m} \sim \text{Po}\left(\sum_{k=1}^K \theta_{n,k} \phi_{m,k}\right) &\iff x_{n,m} = \sum_{k=1}^K z_{n,m,k} \\ z_{n,m,k} &\sim \text{Po}(\theta_{n,k} \phi_{m,k}) \quad \text{independently.}\end{aligned}$$

From this perspective, a user's rating of an item is a sum of ratings along each feature dimension, and each feature rating is an independent Poisson random variable.

The joint distribution is,

$$\begin{aligned}p(\mathbf{X}, \mathbf{Z}, \Theta, \Phi) &= \left[\prod_{n=1}^N \prod_{m=1}^M \mathbb{I}\left[x_{n,m} = \sum_{k=1}^K z_{n,m,k}\right] \prod_{k=1}^K \text{Po}(z_{n,m,k} \mid \theta_{n,k} \phi_{m,k}) \right] \\ &\times \left[\prod_{n=1}^N \prod_{k=1}^K \text{Ga}(\theta_{n,k} \mid \alpha_\theta, \beta_\theta) \right] \times \left[\prod_{m=1}^M \prod_{k=1}^K \text{Ga}(\phi_{m,k} \mid \alpha_\phi, \beta_\phi) \right]\end{aligned}$$

where $\mathbf{Z} \in \mathbb{N}^{N \times M \times K}$ denotes the *tensor* of feature ratings, $\Theta \in \mathbb{R}_+^{N \times K}$ is a matrix with rows θ_n , and $\Phi \in \mathbb{R}_+^{M \times K}$ is a matrix with rows ϕ_m .

Setup

```
In [2]: import torch
from torch.distributions import Distribution, Gamma, Poisson, Multinomial
from torch.distributions.kl import kl_divergence

from tqdm.auto import trange

import matplotlib.pyplot as plt
import seaborn as sns
sns.set_context("notebook")
```

Problem 1: Conditional distributions [math]

Since this model is constructed from conjugate exponential family distributions, the conditionals are available in closed form. We will let $\mathbf{z}_{n,m} = (z_{n,m,1}, \dots, z_{n,m,K})$.

Problem 1a: Derive the conditional for $\mathbf{z}_{n,m}$

Find the conditional density $p(\mathbf{z}_{n,m} \mid x_{n,m}, \boldsymbol{\theta}_n, \boldsymbol{\phi}_m)$.

Your answer here.

$$p(\mathbf{z}_{n,m} \mid x_{n,m}, \boldsymbol{\theta}_n, \boldsymbol{\phi}_m) = \text{Multinomial}(\mathbf{z}_{n,m}; \sum_{k=1}^K z_{n,m,k}, \frac{\theta_{n,1}\phi_{m,1}}{\lambda_{n,m}}, \dots, \frac{\theta_{n,K}\phi_{m,K}}{\lambda_{n,m}})$$

$$\text{where } \lambda_{n,m} = \sum_{k=1}^K \theta_{n,k} \phi_{m,k}$$

Problem 1b: Derive the conditional for $\theta_{n,k}$

Find the conditional density $p(\theta_{n,k} \mid \mathbf{Z}, \boldsymbol{\Phi})$.

Your answer here.

$$p(\theta_{n,k} \mid \mathbf{Z}, \boldsymbol{\Phi}) \propto \text{Gamma}(\theta_{n,k} \mid \alpha_\theta, \beta_\theta) \prod_{m=1}^M \text{Po}(Z_{n,m,k} \mid \theta_{n,k} \phi_{m,k}) \propto \theta_{n,k}^{\alpha_\theta + \sum_{m=1}^M Z_{n,m,k}} e^{(-\theta_{n,k})(\beta_\theta + \sum_{m=1}^M \phi_{m,k})}$$

So the posterior distribution is $\text{Gamma}(\alpha_\theta + \sum_{m=1}^M Z_{n,m,k}, \beta_\theta + \sum_{m=1}^M \phi_{m,k})$

Problem 1c: Derive the conditional for $\phi_{m,k}$

Find the conditional density $p(\phi_{m,k} \mid \mathbf{Z}, \boldsymbol{\Theta})$.

Your answer here.

$$p(\phi_{m,k} \mid \mathbf{Z}, \boldsymbol{\Phi}) \propto \text{Gamma}(\phi_{m,k} \mid \alpha_\phi, \beta_\phi) \prod_{n=1}^N \text{Po}(Z_{n,m,k} \mid \theta_{n,k} \phi_{m,k}) \propto \phi_{m,k}^{\alpha_\phi + \sum_{n=1}^N Z_{n,m,k}} e^{(-\phi_{m,k})(\beta_\phi + \sum_{n=1}^N \theta_{n,k})}$$

So the posterior distribution is $\text{Gamma}(\alpha_\phi + \sum_{n=1}^N Z_{n,m,k}, \beta_\phi + \sum_{n=1}^N \theta_{n,k})$

Problem 2: Coordinate ascent variational inference [math]

We will perform inference in this model using a mean-field variational posterior which factorizes according to:

$$\begin{aligned} q(\mathbf{Z}, \boldsymbol{\Phi}, \boldsymbol{\Theta}) &= q(\mathbf{Z})q(\boldsymbol{\Phi})q(\boldsymbol{\Theta}) \\ &= \left[\prod_{n=1}^N \prod_{m=1}^M q(\mathbf{z}_{n,m}) \right] \left[\prod_{n=1}^N \prod_{k=1}^K q(\theta_{n,k}) \right] \left[\prod_{m=1}^M \prod_{k=1}^K q(\phi_{m,k}) \right] \end{aligned}$$

The optimal mean field factors will have the same forms as the conditional distributions above.

Problem 2a: Derive the CAVI update for $q(\mathbf{z}_{n,m})$

Show that, fixing $q(\Phi)$ and $q(\Theta)$, the optimal $q(\mathbf{z}_{n,m})$ is given by:

$$q(\mathbf{z}_{n,m}; \lambda_{n,m}^{(z)}) = \text{Mult}(\mathbf{z}_{n,m}; x_{n,m}, \lambda_{n,m}^{(z)})$$

$$\log \lambda_{n,m,k}^{(z)} = \mathbb{E}_q[\log \theta_{n,k} + \log \phi_{m,k}] + c$$

Your answer here.

since

$$p(\mathbf{z}_{n,m} | x_{n,m}, \lambda_{n,m}) = \text{Multinomial}(\mathbf{z}_{n,m}; \sum_{k=1}^K z_{n,m,k}, \frac{\theta_{n,1} \phi_{m,1}}{\lambda_{n,m}}, \dots, \frac{\theta_{n,K} \phi_{m,K}}{\lambda_{n,m}}) = \frac{x_{n,m}!}{\prod_{k=1}^K z_{n,m,k}!} e^{\sum_{k=1}^K z_{n,m,k} \log \frac{\theta_{n,k} \phi_{m,k}}{\lambda_{n,m}}}$$

This is an exponential family, hence the optimal CAVI update to each coordinate of the natural parameter is $\mathbb{E}_q[\log \theta_{n,k} + \log \phi_{m,k}] + c$, where $c = -\mathbb{E}_q \log \lambda_{n,m}$. This is the log of the probabilities in the posterior multinomial distribution.

Problem 2b: Derive the CAVI update for $q(\theta_{n,k})$

Show that, fixing $q(\mathbf{Z})$ and $q(\Phi)$, the optimal $q(\theta_{n,k})$ is given by:

$$q(\theta_{n,k}; \lambda_{n,k,1}^{(\theta)}, \lambda_{n,k,2}^{(\theta)}) = \text{Ga}(\theta_{n,k}; \lambda_{n,k,1}^{(\theta)}, \lambda_{n,k,2}^{(\theta)})$$

$$\lambda_{n,k,1}^{(\theta)} = \alpha_{\theta} + \sum_{m=1}^M \mathbb{E}_q[z_{n,m,k}]$$

$$\lambda_{n,k,2}^{(\theta)} = \beta_{\theta} + \sum_{m=1}^M \mathbb{E}_q[\phi_{m,k}]$$

Your answer here.

Gamma is one of the exponential family, and both the prior and the posterior of $\theta_{n,k}$ is the Gamma distribution, the optimal CAVI update involves taking expectation of the posterior

parameters. Since the true posterior is $\text{Gamma}(\alpha_\theta + \sum_{m=1}^M Z_{n,m,k}, \beta_\theta + \sum_{m=1}^M \phi_{m,k})$. The optimal CAVI update is $\text{Gamma}(\alpha_\theta + \sum_{m=1}^M \mathbb{E}_q[Z_{n,m,k}], \beta_\theta + \sum_{m=1}^M \mathbb{E}_q[\phi_{m,k}])$

Problem 2c: Derive the CAVI update for $q(\phi_{m,k})$

Show that, fixing $q(\mathbf{Z})$ and $q(\Theta)$, the optimal $q(\phi_{m,k})$ is given by:

$$\begin{aligned} q(\phi_{m,k}; \lambda_{m,k,1}^{(\phi)}, \lambda_{m,k,2}^{(\phi)}) &= \text{Ga}(\phi_{m,k}; \lambda_{m,k,1}^{(\phi)}, \lambda_{m,k,2}^{(\phi)}) \\ \lambda_{m,k,1}^{(\phi)} &= \alpha_\phi + \sum_{n=1}^N \mathbb{E}_q[z_{n,m,k}] \\ \lambda_{m,k,2}^{(\phi)} &= \beta_\phi + \sum_{n=1}^N \mathbb{E}_q[\theta_{n,k}] \end{aligned}$$

Your answer here.

Gamma is one of the exponential family, and both the prior and the posterior of $\phi_{m,k}$ is the Gamma distribution, the optimal CAVI update involves taking expectation of the posterior parameters. Since the true posterior is $\text{Gamma}(\alpha_\phi + \sum_{n=1}^N Z_{n,m,k}, \beta_\phi + \sum_{n=1}^N \theta_{n,k})$. The optimal CAVI update is $\text{Gamma}(\alpha_\phi + \sum_{n=1}^N \mathbb{E}_q[Z_{n,m,k}], \beta_\phi + \sum_{n=1}^N \mathbb{E}_q[\theta_{n,k}])$

Problem 2d: Find the expected sufficient statistics

To update the variational factors, we need the expectations $\mathbb{E}_q[z_{n,m,k}]$, $\mathbb{E}_q[\log \theta_{n,k} + \log \phi_{m,k}]$, $\mathbb{E}_q[\theta_{n,k}]$, and $\mathbb{E}_q[\phi_{m,k}]$. Assume that each factor follows the forms derived above. That is, assume $q(\mathbf{z}_{n,m})$ is multinomial with parameters $\lambda_{n,m}^{(z)}$ while $q(\theta_{n,k})$ and $q(\phi_{m,k})$ are gamma with parameters $(\lambda_{n,k,1}^{(\theta)}, \lambda_{n,k,2}^{(\theta)})$ and $(\lambda_{m,k,1}^{(\phi)}, \lambda_{m,k,2}^{(\phi)})$, respectively. Derive what each of these expectations are in closed form.

Your answer here.

$$\mathbb{E}_q[z_{n,m,k}] = x_{n,m} \lambda_{n,m,k}^{(z)}$$

$$\mathbb{E}_q[\theta_{n,k}] = \frac{\lambda_{n,k,1}^{(\theta)}}{\lambda_{n,k,2}^{(\theta)}}$$

$$\mathbb{E}_q[\phi_{m,k}] = \frac{\lambda_{m,k,1}^{(\phi)}}{\lambda_{m,k,2}^{(\phi)}}$$

Problem 3: Implement Coordinate Ascent Variational Inference [code]

First we'll give some helper functions and objects. Because PyTorch doesn't offer support for batched multinomial distributions in which the total counts differ (e.g. each $\mathbf{z}_{n,m}$ follows a multinomial distribution in which the total count is $x_{n,m}$), we have defined a

`BatchedMultinomial` distribution for your convenience. This distribution doesn't support sampling, but will return the mean of each Multinomial variable in its batch. This is exactly what is needed for the CAVI updates.

```
In [3]: def gamma_expected_log(gamma_distbn):
        """Helper function to compute the expectation of log(X) where X follows a
        gamma distribution.
        """
        return torch.digamma(gamma_distbn.concentration) - torch.log(gamma_distbn.rate)

class BatchedMultinomial(Multinomial):
    """
    Creates a Multinomial distribution parameterized by `total_count` and
    either `probs` or `logits` (but not both). The innermost dimension of
    `probs` indexes over categories. All other dimensions index over batches.

    The `probs` argument must be non-negative, finite and have a non-zero sum,
    and it will be normalized to sum to 1 along the last dimension. `probs` will
    return this normalized value. The `logits` argument will be interpreted as
    unnormalized log probabilities and can therefore be any real number. It will
    likewise be normalized so that the resulting probabilities sum to 1 along
    the last dimension. `logits` will return this normalized value.

    Args:
        total_count (Tensor): number of trials
        probs (Tensor): event probabilities
            Has shape total_count.shape + (num_categories,)
        logits (Tensor): event log probabilities (unnormalized)
            Has shape total_count.shape + (num_categories,)

    Note: this text is mostly from the PyTorch documentation for the
        Multinomial distribution
    """
    def __init__(self, total_count, probs=None, logits=None, validate_args=None):
        super().__init__(probs=probs, logits=logits, validate_args=validate_args)
        self.total_count = total_count

    @property
    def mean(self):
        return self.total_count[..., None] * self.probs
```

Problem 3a: Implement a CAVI update step

Using the update equations derived in Problem 2, complete the `cavi_step` function below.

Hint: Given a `Distribution` named `d`, `d.mean` returns the mean of that distribution.

```
In [28]: import torch.nn.functional as F

def cavi_step(X, q_z, q_theta, q_phi, alpha_theta, beta_theta, alpha_phi, beta_phi)
    """One step of CAVI.

    Args:
        X: torch.tensor of shape (N, M)
        q_z: variational posterior over z, BatchedMultinomial distribution
        q_theta: variational posterior over theta, Gamma distribution
        q_phi: variational posterior over eta, Gamma distribution

    Returns:
        (q_z, q_theta, q_phi): Updated distributions after performing CAVI updates
    """
    ###
    # Your code here
    N, M = X.shape
    K = q_theta.concentration.shape[1] # Latent dimension

    log_theta = gamma_expected_log(q_theta)
    log_phi = gamma_expected_log(q_phi)

    log_theta_expand = log_theta[:, None, :]
    log_phi_expand = log_phi[None, :, :]
    log_lambda = log_theta_expand + log_phi_expand

    probs = F.softmax(log_lambda, dim=-1)
    q_z = BatchedMultinomial(total_count=X, probs=probs)

    expected_z = q_z.mean
    sum_z_over_m = expected_z.sum(dim=1)
    expected_phi = q_phi.mean
    sum_phi_over_m = expected_phi.sum(dim=0).unsqueeze(0).expand(N, K)

    q_theta = Gamma(
        concentration=alpha_theta + sum_z_over_m,
        rate=beta_theta + sum_phi_over_m
    )

    sum_z_over_n = expected_z.sum(dim=0)
    expected_theta = q_theta.mean
    sum_theta_over_n = expected_theta.sum(dim=0).unsqueeze(0).expand(M, K)

    q_phi = Gamma(
        concentration=alpha_phi + sum_z_over_n,
        rate=beta_phi + sum_theta_over_n
    )
```

```
return q_z, q_theta, q_phi
```

Problem 3b: ELBO Calculation [math]

Recall that the evidence lower bound is defined as:

$$\mathcal{L}(q) = \mathbb{E}_q [\log p(\mathbf{X}, \mathbf{Z}, \Phi, \Theta) - \log q(\mathbf{Z}, \Phi, \Theta)]$$

Assume that $q(\mathbf{Z})$ has support contained in $\{\mathbf{Z} : \sum_{k=1}^K z_{n,m,k} = x_{n,m} \text{ for all } n, m\}$. Show that we can rewrite $\mathcal{L}(q)$ as:

$$\mathcal{L}(q) = \mathbb{E}_q [\log p(\mathbf{Z} \mid \Theta, \Phi) - \log q(\mathbf{Z})] - \text{KL}(q(\Theta) \parallel p(\Theta)) - \text{KL}(q(\Phi) \parallel p(\Phi))$$

Next, use that $q(\mathbf{z}_{n,m}; \lambda_{n,m}^{(z)}) = \text{Mult}(\mathbf{z}_{n,m}; x_{n,m}, \lambda_{n,m}^{(z)})$ and by plug in the densities of the Poisson and Multinomial distributions to show that we have:

$$\begin{aligned} \mathbb{E}_q [\log p(\mathbf{Z} \mid \Theta, \Phi) - \log q(\mathbf{Z})] = \\ \sum_{n=1}^N \sum_{m=1}^M \mathbb{E}_q \left[\sum_{k=1}^K -\theta_{n,k} \phi_{m,k} + z_{n,m,k} \log(\theta_{n,k} \phi_{m,k}) - z_{n,m,k} \log(\lambda_{n,m,k}^{(z)}) \right] - \log(x_{n,m}!) \end{aligned}$$

Explain why we have:

$$\begin{aligned} \mathbb{E}_q \left[-\theta_{n,k} \phi_{m,k} + z_{n,m,k} \log(\theta_{n,k} \phi_{m,k}) - z_{n,m,k} \log(\lambda_{n,m,k}^{(z)}) \right] = \\ - \mathbb{E}_q [\theta_{n,k}] \mathbb{E}_q [\phi_{m,k}] + \mathbb{E}_q [z_{n,m,k}] \left(\mathbb{E}_q [\log(\theta_{n,k})] + \mathbb{E}_q [\log(\phi_{m,k})] - \log(\lambda_{n,m,k}^{(z)}) \right) \end{aligned}$$

Your answer here.

$\mathcal{L}(q) = \mathbb{E}_q [\log p(\mathbf{X}, \mathbf{Z}, \Phi, \Theta) - \log q(\mathbf{Z}, \Phi, \Theta)] = \mathbb{E}_q [\log p(\mathbf{X}, \mathbf{Z} \mid \Phi, \Theta) + \log p(\Phi, \Theta) - \log q(\mathbf{Z}, \Phi, \Theta)]$ by the mean-field assumption.

Since $\log p(\mathbf{X}, \mathbf{Z} \mid \Phi, \Theta) = \log p(\mathbf{Z} \mid \Phi, \Theta) \mathbf{1}_{\sum_{k=1}^K z_{n,m,k} = x_{n,m} \forall n, m'}$ we have

$$\mathcal{L}(q) = \mathbb{E}_q [\log p(\mathbf{X}, \mathbf{Z}, \Phi, \Theta) - \log q(\mathbf{Z}, \Phi, \Theta)] = \mathbb{E}_q [\log p(\mathbf{Z} \mid \Phi, \Theta) + \log p(\Phi) + \log p(\Theta) - \text{KL}(q(\Theta) \parallel p(\Theta)) - \text{KL}(q(\Phi) \parallel p(\Phi))]$$

$$\begin{aligned} \mathbb{E}_q [\log p(\mathbf{Z} \mid \Theta, \Phi) - \log q(\mathbf{Z})] &= \sum_{n=1}^N \sum_{m=1}^M \mathbb{E}_q [\sum_{k=1}^K (\log p(Z_{n,m,k} \mid \Theta, \Phi)) - \log q(\mathbf{Z}_{n,m})] \\ &= \sum_{m=1}^M \mathbb{E}_q [\sum_{k=1}^K (Z_{n,m,k} \log(\theta_{n,k} \phi_{m,k}) - \theta_{n,k} \phi_{m,k} - \log Z_{n,m,k}!) - \log(X_{n,m}! + \sum_{k=1}^K \log Z_{n,m,k})] \\ &= \sum_{m=1}^M \mathbb{E}_q \left[\sum_{k=1}^K -\theta_{n,k} \phi_{m,k} + z_{n,m,k} \log(\theta_{n,k} \phi_{m,k}) - z_{n,m,k} \log(\lambda_{n,m,k}^{(z)}) \right] - \log(x_{n,m}!) \end{aligned}$$

By the mean-field assumption $\theta_{n,k}, \phi_{m,k}, z_{n,m,k}$ are independent, thus we have

$$\begin{aligned} \mathbb{E}_q \left[-\theta_{n,k} \phi_{m,k} + z_{n,m,k} \log(\theta_{n,k} \phi_{m,k}) - z_{n,m,k} \log(\lambda_{n,m,k}^{(z)}) \right] = \\ - \mathbb{E}_q [\theta_{n,k}] \mathbb{E}_q [\phi_{m,k}] + \mathbb{E}_q [z_{n,m,k}] \left(\mathbb{E}_q [\log(\theta_{n,k})] + \mathbb{E}_q [\log(\phi_{m,k})] - \log(\lambda_{n,m,k}^{(z)}) \right) \end{aligned}$$

Problem 3c: Implement the ELBO [code]

Using our expression above, write a function which evaluates the evidence lower bound.

Hints:

- Use the `kl_divergence` function imported above to compute the KL divergence between two `Distributions` in the same family.
- Recall that for integers n , $\Gamma(n + 1) = n!$ where Γ is the [Gamma function](#). $\log \Gamma$ is implemented in PyTorch as `torch.lgamma`.

```
In [8]: def elbo(X, q_z, q_theta, q_phi, p_theta, p_phi):
        """Compute the evidence lower bound.

        Args:
            X: torch.tensor of shape (N, M)
            q_z: variational posterior over z, BatchedMultinomial distribution
            q_theta: variational posterior over theta, Gamma distribution
            q_phi: variational posterior over eta, Gamma distribution
            p_theta: prior over theta, Gamma distribution
            p_phi: prior over eta, Gamma distribution

        Returns:
            elbo: torch.tensor of shape []
        """
        N, M = X.shape
        K = q_theta.concentration.shape[1]
        E_log_theta = gamma_expected_log(q_theta) # (N, K)
        E_log_phi = gamma_expected_log(q_phi) # (M, K)
        E_theta = q_theta.mean # (N, K)
        E_phi = q_phi.mean # (M, K)
        E_z = q_z.mean # (N, M, K)
        log_lambda = torch.log(q_z.probs + 1e-12) # (N, M, K)

        E_log_theta_phi = E_log_theta[:, None, :] + E_log_phi[None, :, :] # (N, M, K)
        E_theta_phi = torch.einsum("nk,mk->nmk", E_theta, E_phi) # (N, M, K)

        term1 = torch.sum(E_z * E_log_theta_phi) - torch.sum(E_theta_phi)
        term2 = -torch.sum(E_z * log_lambda)
        term3 = -torch.sum(torch.lgamma(X + 1))

        kl_theta = kl_divergence(q_theta, p_theta).sum()
        kl_phi = kl_divergence(q_phi, p_phi).sum()

        elbo = term1 + term2 + term3 - kl_theta - kl_phi
        return elbo / torch.sum(X)
```

Implement CAVI loop [given]

Using your functions defined above, complete the function `cavi` below. `cavi` loops for some number of iterations, updating each of the variational factors in sequence and

evaluating the ELBO at each step.

```
In [29]: from torch.distributions import Uniform

def cavi(data,
        num_factors=10,
        num_iters=100,
        tol=1e-5,
        alpha_theta=0.1,
        beta_theta=1.0,
        alpha_phi=0.1,
        beta_phi=1.0,
        seed=0
    ):
    """Run coordinate ascent VI for Poisson matrix factorization.

    Args:

    Returns:
        elbos, (q_z, q_theta, q_phi):
    """
    data = data.float()
    N, M = data.shape
    K = num_factors    # short hand

    # Initialize the variational posteriors.
    q_phi = Gamma(Uniform(0.5 * alpha_phi, 1.5 * alpha_phi).sample((M, K)),
                  Uniform(0.5 * beta_phi, 1.5 * beta_phi).sample((M, K)))
    q_theta = Gamma(Uniform(0.5 * alpha_theta, 1.5 * alpha_theta).sample((N, K)),
                   Uniform(0.5 * beta_theta, 1.5 * beta_theta).sample((N, K)))
    q_z = BatchedMultinomial(data, logits=torch.zeros((N, M, K)))

    p_theta = Gamma(alpha_theta, beta_theta)
    p_phi = Gamma(alpha_phi, beta_phi)

    # Run CAVI
    elbos = [elbo(data, q_z, q_theta, q_phi, p_theta, p_phi)]
    for itr in trange(num_iters):
        q_z, q_theta, q_phi = cavi_step(data, q_z, q_theta, q_phi,
                                         alpha_theta, beta_theta,
                                         alpha_phi, beta_phi)

        elbos.append(elbo(data, q_z, q_theta, q_phi, p_theta, p_phi))
    return torch.tensor(elbos), (q_z, q_theta, q_phi)
```

Test your implementation on a toy dataset

To check your implementation is working properly, we will fit a mean-field variational posterior using data sampled from the true model.

```
In [30]: # Constants
N = 100    # num "users"
M = 1000   # num "items"
```

```

K = 5      # number of Latent factors

# Hyperparameters
alpha = 0.1 # sparse gamma prior with mean alpha/beta
beta = 1.0

# Sample data from the model
torch.manual_seed(305)
theta = Gamma(alpha, beta).sample(sample_shape=(N, K))
phi = Gamma(alpha, beta).sample(sample_shape=(M, K))
data = Poisson(theta @ phi.T).sample()

print(data.shape)
# Plot the data matrix
plt.imshow(data, aspect="auto", vmax=5, cmap="Greys")
plt.xlabel("items")
plt.ylabel("users")
plt.colorbar()


print("Max data: ", data.max())
print("num zeros: ", torch.sum(data == 0))

```

```

torch.Size([100, 1000])
Max data:  tensor(14.)
num zeros: tensor(95568)

```




```

In [31]: elbos, (q_z, q_theta, q_phi) = cavi(data)
0%|          | 0/100 [00:00<?, ?it/s]

```

```
In [32]: plt.plot(elbos[1:])
plt.xlabel("Iteration")
plt.ylabel("ELBO per entry")
```

```
Out[32]: Text(0, 0.5, 'ELBO per entry')
```




```
In [33]: true_rates = theta @ phi.T
inf_rates = q_theta.mean @ q_phi.mean.T

# Plot the data matrix
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
plt.imshow(true_rates, aspect="auto", vmax=3, cmap="Greys")
plt.xlabel("items")
plt.ylabel("users")
plt.title("true rates")
plt.colorbar()

plt.subplot(1, 2, 2)
plt.imshow(inf_rates, aspect="auto", vmax=3, cmap="Greys")
plt.xlabel("items")
plt.ylabel("users")
plt.title("inferred rates")
plt.colorbar()
```

```
Out[33]: <matplotlib.colorbar.Colorbar at 0x20d0601b5f0>
```



Problem 4: Run your code on a downsampled LastFM dataset

Next, we will use data gathered from [Last.FM](#) users to fit a PMF model. We use a downsampled version of the [Last.FM-360K users](#) dataset. This dataset records how many times each user played an artist's songs. We downsample the data to include only the 2000 most popular artists, as measured by how many users listened to the artist at least once, and the 1000 most prolific users, as measured by how many artists they have listened to.

In the code below, we use `lfm` to represent the data matrix X in the model. That is, `lfm[n, d]` denotes how many times the `n`-th user played a song by the `d`-th artist.

```
In [34]: import pandas as pd

lfm_df = pd.read_csv('subsampled_last_fm.csv')
lfm = lfm_df.pivot_table(index='UserID', columns='ItemID', values='Count', aggfunc=
    .fillna(0).astype(int).to_numpy())
lfm = torch.tensor(lfm, dtype=torch.int)
print(lfm.shape)
```

```
torch.Size([999, 2000])
```


C:\Users\HaoLi\AppData\Local\Temp\ipykernel_176068\1531818262.py:4: FutureWarning: The provided callable <built-in function sum> is currently using DataFrameGroupBy.sum. In a future version of pandas, the provided callable will be used directly. To keep current behavior pass the string "sum" instead.

```
lfm = lfm_df.pivot_table(index='UserID', columns='ItemID', values='Count', aggfunc=
    sum)\
```

```
In [35]: plt.imshow(lfm, aspect="auto", vmax=100, cmap="Greys")
plt.xlabel("items")
```

```
plt.ylabel("users")
plt.colorbar()
```

Out[35]: <matplotlib.colorbar.Colorbar at 0x20d0422f0b0>



Using the code below, run coordinate ascent variational inference on this dataset. Our implementation takes around 10-15 minutes to finish, and achieves a rescaled ELBO of around -2 .

```
In [36]: elbos, (q_z, q_theta, q_phi) = cavi(lfm,
        num_factors=40,
        num_iters=200,
        alpha_theta=1.,
        beta_theta=0.5,
        alpha_phi=1.,
        beta_phi=0.5)
```

0%| | 0/200 [00:00<?, ?it/s]

Investigate "genres"

The columns of \mathbf{H} correspond to weights on artists. Intuitively, each of the K columns should put weight on subsets of artists that are often played together. We might think of these columns as reflecting different "genres" of music. The code below the top 10 artists for a few of these columns.

```
In [38]: # Find the 10 most used genres
genre_loading = q_theta.mean.sum(0)
genre_order = torch.argsort(genre_loading, descending=True)

# Print the top 10 artists for each of the top 10 genres
for genre in genre_order[:10]:
    print("genre ", genre)
    artist_idx = torch.argsort(q_phi.mean[:, genre],
                              descending=True)[:10].numpy()
    subset = lfm_df[lfm_df['ItemID'].isin(artist_idx)]
    print(subset[['ItemID', 'Artist']].drop_duplicates())
    print("")
```

genre tensor(1)

	ItemID	Artist
1	1647	tegan and sara
10	15	sufjan stevens
24	38	cat power
44	28	elliott smith
154	1306	modest mouse
155	1578	ryan adams
157	910	bright eyes
469	1959	fleet foxes
1019	815	sam cooke
1289	1933	jay-z

genre tensor(7)

	ItemID	Artist
49	1286	radiohead
204	627	david bowie
206	795	the cure
320	1658	devo
394	476	the smiths
475	12	morrissey
687	1877	new order
711	692	sonic youth
739	813	guided by voices
757	469	the magnetic fields

genre tensor(25)

	ItemID	Artist
60	990	pink floyd
61	758	metallica
105	773	led zeppelin
142	1591	iron maiden
165	806	judas priest
211	763	ac/dc
1365	614	black sabbath
1369	1301	megadeth
1375	651	motörhead
1677	1768	kiss
20257	763	acdc

genre tensor(14)

	ItemID	Artist
55	44	jimi hendrix
56	1377	the beatles
60	990	pink floyd
69	857	bob dylan
73	1210	the doors
105	773	led zeppelin
204	627	david bowie
551	988	pearl jam
563	1100	the clash
2920	1855	guns n' roses
2934	1855	guns n roses
18201	1377	beatles

genre tensor(29)

	ItemID	Artist
336	442	ghostface
817	1150	a tribe called quest
1419	1500	common
1431	964	the roots
2297	1639	nas
2312	1213	mos def
2693	1598	j dilla
3613	1834	madlib
3835	99	wu-tang clan
4726	442	ghostface killah
10000	99	wu tang clan
15054	1319	blessthefall
24609	964	the roots featuring d'angelo

genre tensor(2)

	ItemID	Artist
13	676	stars
38	1605	coldplay
49	1286	radiohead
83	1196	muse
109	981	jason mraz
167	2	death cab for cutie
235	822	m83
282	1066	bloc party
1080	1785	incubus
1852	1279	snow patrol
29203	676	the stars

genre tensor(32)

	ItemID	Artist
56	1377	the beatles
67	26	queen
106	730	eric clapton
116	63	frédéric chopin
183	911	madonna
280	1705	abba
324	1370	the rolling stones
330	13	elvis presley
565	1254	u2
876	832	mushroomhead
18201	1377	beatles

genre tensor(28)

	ItemID	Artist
170	37	daft punk
218	1007	depeche mode
222	574	blondie
865	1494	the knife
951	1930	fever ray
973	1379	crystal castles
992	657	yeah yeah yeahs
1303	1400	ladytron
1443	1313	franz ferdinand
1447	1947	justice

genre	tensor(5)	
	ItemID	Artist
59	645	miles davis
64	1672	johnny cash
69	857	bob dylan
277	1043	various artists
293	1533	tom waits
297	163	nick cave and the bad seeds
324	1370	the rolling stones
330	13	elvis presley
343	755	leonard cohen
443	163	nick cave & the bad seeds
913	848	bruce springsteen
18256	1043	v.a.

genre	tensor(0)	
	ItemID	Artist
328	1819	bad brains
452	1791	have heart
454	820	comeback kid
788	1692	ramones
821	277	rancid
860	1735	nofx
885	1281	against me!
894	142	bad religion
1364	1170	black flag
11243	1769	sick of it all

Problem 4a

Inspect the data either using the csv file or the pandas dataframe and choose a user who has listened to artists you recognize. If you are not familiar with any of the artists, use the listener with UserID 349, who mostly listens to hip-hop artists. For the particular user n you choose, find the 10 artists who are predicted to have the most plays by sorting the vector of mean song counts predicted by the model, i.e. the n^{th} row of $\mathbb{E}_q[\Theta\Phi^T]$. Are these artists you would expect the user would enjoy? Are there any artists that the user has not listened to?

Hint: Use `torch.argsort(..., descending=True)` to return the indices of the largest elements of a vector in descending order.

```
In [50]: ###
artist_idx=torch.argsort((q_theta.mean@q_phi.mean.T)[736], descending=True)[:10].nu
subset = lfm_df[lfm_df['ItemID'].isin(artist_idx)]
print(subset[['ItemID', 'Artist']].drop_duplicates())
print("")
##
```

	ItemID	Artist
49	1286	radiohead
204	627	david bowie
206	795	the cure
320	1658	devo
394	476	the smiths
475	12	morrissey
687	1877	new order
711	692	sonic youth
739	813	guided by voices
757	469	the magnetic fields

UserID 736 listens to many Japanese musicians as can be seen from the original .csv file. The predicted list of top artists that he would listen to does not seem plausible.

Problem 5: Reflections

Problem 5a

Discuss one advantage and one disadvantage of fitting a posterior using variational inference vs. sampling from the posterior using MCMC.

Your answer here.

VI is typically much faster and more scalable than sampling-based methods like MCMC, especially on large datasets, because VI converts inference into an optimization problem that can be solved efficiently using gradient-based procedures without generating long sample chains.

VI often underestimates the complexity of the true posterior by searching only in a family of fully factorized variational distributions. Also it may not converge to the true posterior because non-convex optimization is involved.

Problem 5b

First, explain why the assumption that \mathbf{Z} , Φ and Θ are independent in the posterior will never hold.

Next, recall that maximizing the ELBO is equivalent to minimizing the KL divergence between the approximate posterior and the true posterior. In general, how will the approximate

posterior differ from the true posterior, given that the variational family does not include the true posterior?

Your answer here.

First, the assumption that \mathbf{Z} , Φ , and Θ are independent in the variational posterior will never hold in general because these variables are inherently dependent in the true posterior. For example, the latent variable \mathbf{Z} depends on both user preferences Θ and item features Φ when generating the observed data. If a user's preference for a topic increases, it can reduce the need for higher weights in item features to explain the same observation. Besides, \mathbf{Z} also depends on both Φ and Θ .

If the variational family does not contain the true posterior, then the approximate posterior will likely be the one in the variational family with the smallest KL divergence w.r.t. the posterior. Moreover, since VI minimizes the KL divergence, the approximate posterior will have very small probabilities where the true posterior has low density. It tends to favor underestimating the posterior variance and capturing only a single mode.

Problem 5c

Suppose we are using this model to recommend new items to users. Describe one improvement that could be made to the model which you think would lead to better recommendations.

Your answer here.

One improvement would be to incorporate side information about users or items—such as user demographics or item metadata (e.g., artist genres, song tags)—into the model. This could be done by embedding the side information as features, and then let Θ and Φ depend on these features to make the variational distributions more informative, or directly let the features influence \mathbf{X} . Incorporating this auxiliary information helps the model make better predictions for new users or sparse users who have interacted with few items.

Submission Instructions

Formatting: check that your code does not exceed 80 characters in line width. If you're working in Colab, you can set *Tools* → *Settings* → *Editor* → *Vertical ruler column* to 80 to see when you've exceeded the limit.

Download your notebook in .ipynb format and use the following commands to convert it to PDF:

```
jupyter nbconvert --to pdf hw5_yourname.ipynb
```

Dependencies:

- `nbconvert` : If you're using Anaconda for package management,

```
conda install -c anaconda nbconvert
```

Upload your .pdf files to Gradescope.